



Fig 4. Resource Footprint across Quantization (Llama-3-8B)


VRAM (GB)



Average power draw (W)



Latency (ms/token)



Throughput (tokens/s)

